Thank you for your submission of Development Individual Project: Presentation.

What was done well:

You selected an appropriate public dataset and provided a clear explanation of its characteristics, limitations and relevance to the comparison. The preprocessing steps and experimental setup were described clearly and supported by a well-documented notebook.

The comparison between the HOG + Linear SVM model and the CNN was structured logically and both approaches were evaluated using the same validation strategy. The rationale for the main design choices was generally explained and supported by evidence from the experiments.

You applied stratified five-fold cross-validation alongside a held-out test set and explained why validation is important for assessing generalisation and reducing the risk of overfitting. The notebook demonstrated careful attention to reproducibility and avoidance of data leakage.

The use of LIME and Integrated Gradients addressed the explainability requirement and the discussion recognised some of the limitations of explanation methods. The robustness analysis also extended the comparison beyond standard performance metrics.

What could be improved:

Much of the presentation focused on reporting results rather than critically evaluating them. You should have discussed alternative modelling approaches and explained more clearly why the selected architectures were preferable to other reasonable choices.

The bias analysis was limited by the dataset and mainly focused on class disparity and image degradation. While this was acknowledged, the discussion remained relatively narrow and could have explored the implications of these limitations in greater depth.

The ethics section identified relevant issues such as privacy, fairness and representation bias, but the discussion of real-world harms and the consequences of classification errors was brief. You should have provided a more detailed consideration of how different stakeholders may be affected.

Although explainability techniques were applied appropriately, the interpretation of the explanations could have been developed further. The discussion often described the outputs without fully examining what they reveal about model behaviour and limitations.

The deployment section provided a reasonable high-level pipeline, but the operational challenges of deployment, monitoring, scalability and retraining were only considered briefly.

General comments for future work:

You addressed the main assessment requirements and demonstrated a good understanding of machine learning model development and evaluation. Future work should focus more on critical analysis and less on describing methods and results. Where conclusions or recommendations are made, they should be supported by a deeper discussion of limitations, trade-offs and practical implications. More attention to fairness, stakeholder impact and deployment challenges would strengthen the overall evaluation.